



Performance Benchmarking of YOLO Architectures for Real-Time Object Detection

A Comparative Study of **YOLOv5**, **YOLOv8**, and **YOLOv11**

Team Members
Beyza GÜLER
Ramazan YILDIZ
Abdelrahman MOHAMED

Department
Computer Engineering
Ataturk University
CUFIV

Agenda

1

Introduction & Problem Statement

Motivation and objectives

3

YOLO Architecture & Selection

Model variants and advantages

5

Methodology & Training Pipeline

Implementation and configuration

7

Model Comparison & Analysis

Benchmarking and insights

9

Challenges & Future Work

Limitations and recommendations

2

Background Concepts

Object detection fundamentals

4

Dataset & Annotation Format

Roboflow fruits detection dataset

6

Experimental Results

YOLOv5, YOLOv8, YOLOv11 performance

8

Deployment & Web Application

Streamlit-based interface



Comprehensive Analysis

Abstract

This project presents a comprehensive **real-time object detection system** based on YOLO architectures, with focus on performance benchmarking across multiple model generations.

Three Model Architectures

Benchmarking YOLOv5, YOLOv8, and YOLOv11 variants

Custom Fine-Tuning

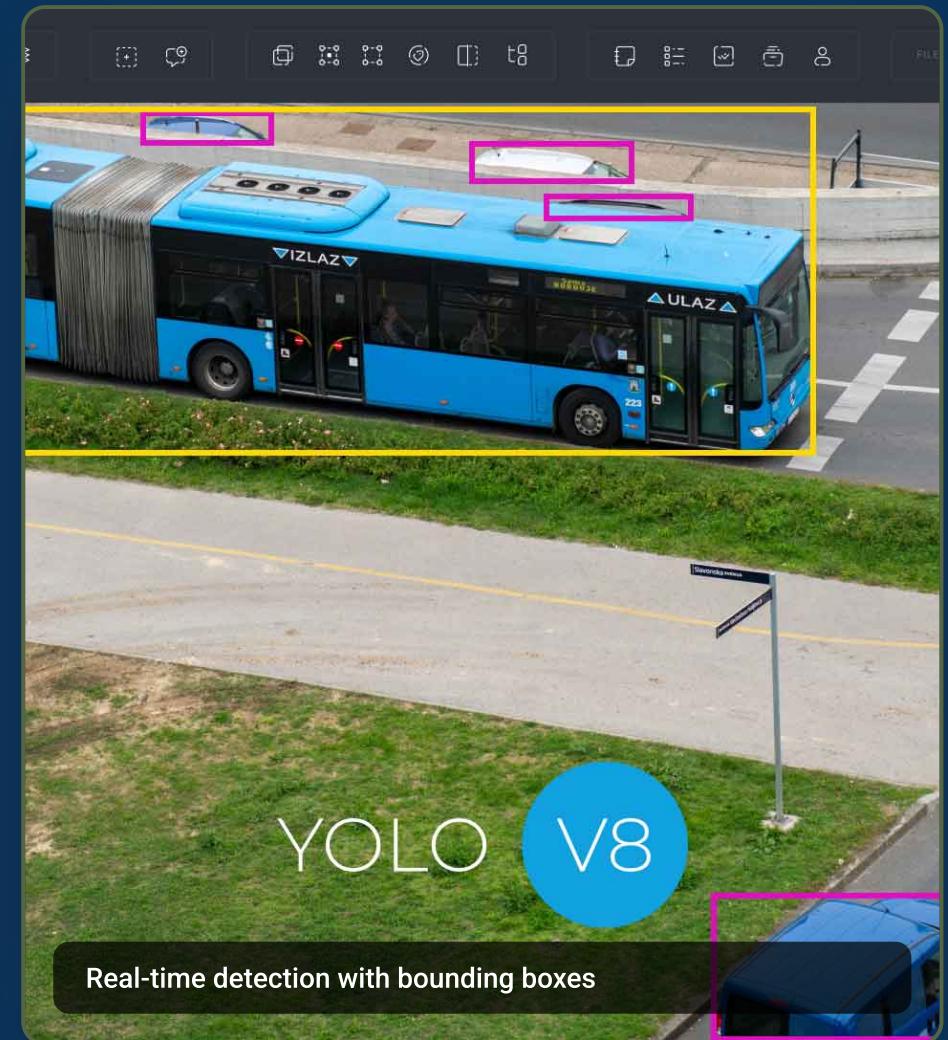
Pre-trained weights adapted to Roboflow fruit dataset

Streamlit Deployment

Interactive web interface for live detection

Key Metrics

Accuracy (mAP), Processing Speed (FPS), Real-time Performance



Problem Statement

The Challenge of Real-Time Object Detection

Speed vs Accuracy Tradeoff

High accuracy models often sacrifice real-time performance

Limited Accessibility

Complex deployment pipelines hinder adoption

Hardware Constraints

GPU availability and computational cost barriers

*"Our goal: Build a balanced system that achieves **~30 FPS** while maintaining high detection accuracy"*

Key Application Areas



Autonomous Driving



Security Systems



Smart Cities

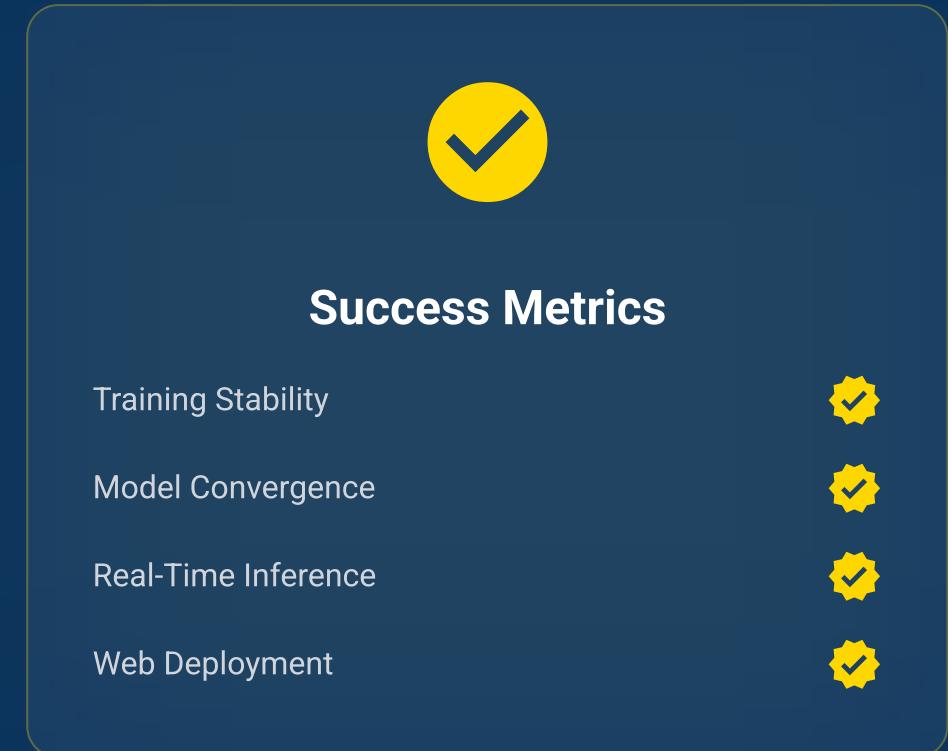
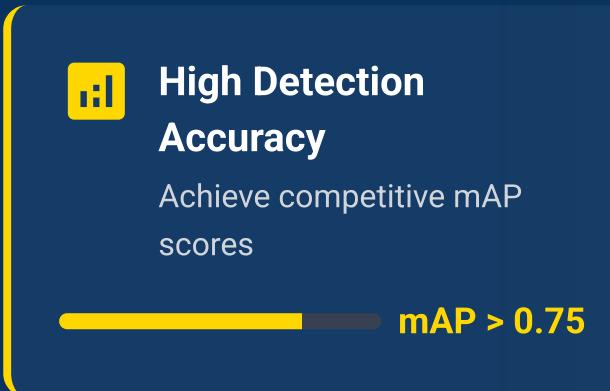
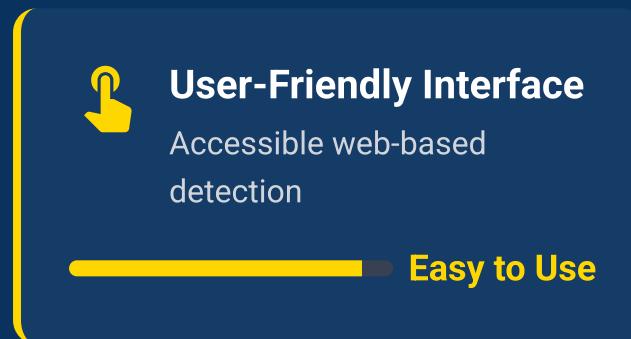
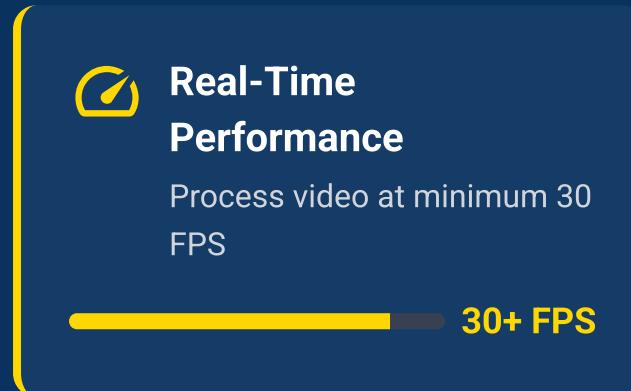


Industrial Automation

 Real-time detection critical for safety & efficiency

Goals & Success Criteria

Project Objectives



Background: Object Detection vs Classification

Image Classification



Single Label Assignment

Assigns ONE category to the entire image

Example: "cat" for an image containing a cat

- ✓ Simpler task
- ✓ Single prediction
- ✓ High computational efficiency

Object Detection



Multiple Objects + Localization

Detects, classifies, and locates ALL objects

Example: "cat (box 1), dog (box 2), person (box 3)"

- ✓ Bounding box regression
- ✓ Multiple predictions
- ✓ Real-time requirement

Why YOLO?

YOLO: You Only Look Once



Single-Stage Detector

Detection in ONE forward pass through the network



Real-Time Performance

Achieves high FPS rates suitable for video streams



Multiple Model Sizes

Nano, Small, Medium, Large variants for different needs



Strong Community Support

Extensive documentation, pre-trained weights, and tutorials



Proven Track Record



Widely Adopted

Industry standard for real-time detection



Accuracy-Speed Balance

Optimal tradeoff for most applications



Easy Integration

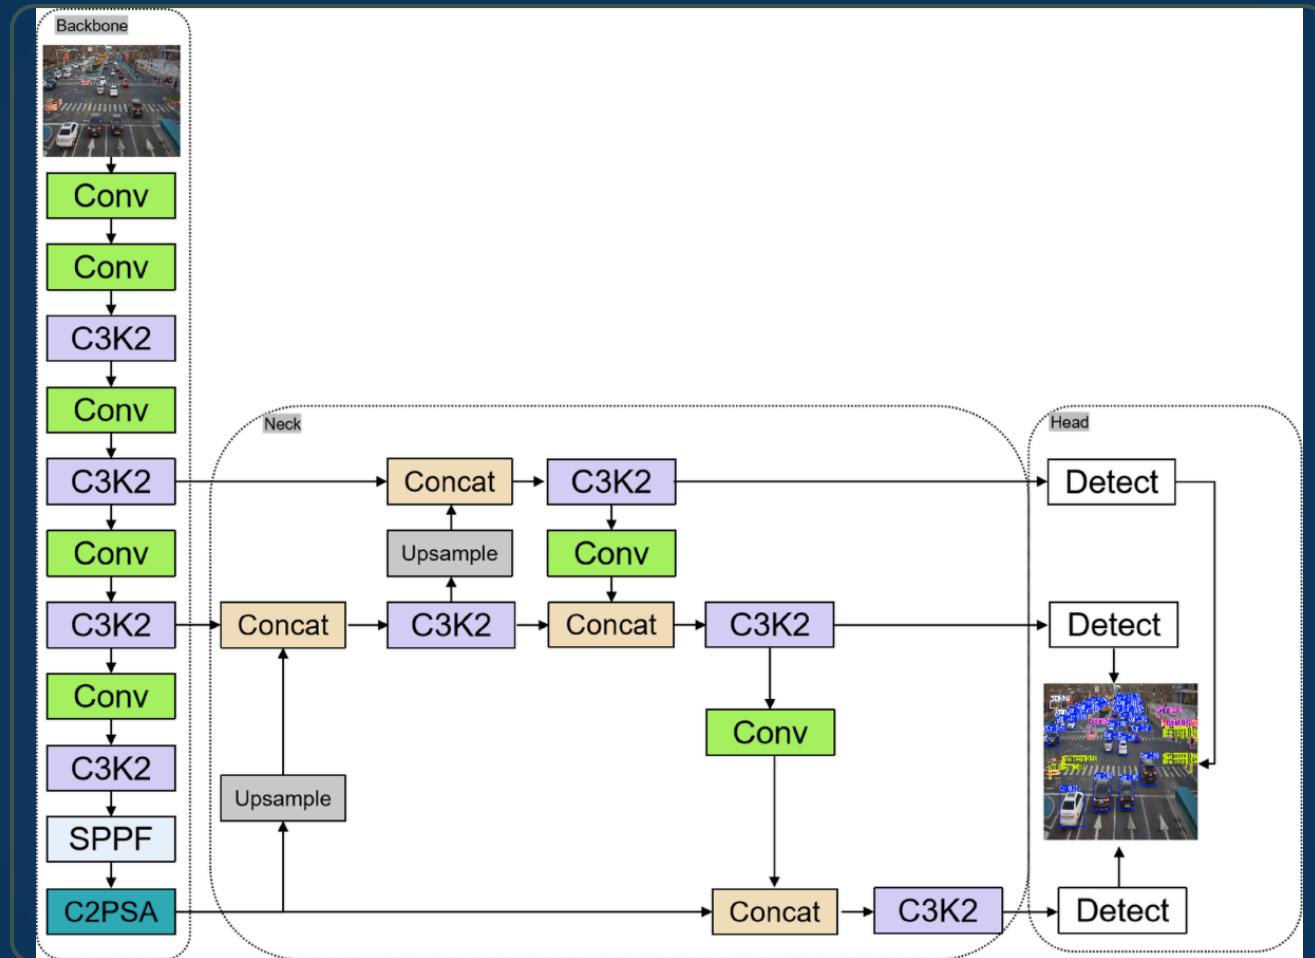
Simple deployment with PyTorch



Battle-Tested Architecture

YOLOv5 → YOLOv8 → YOLOv11 evolution

YOLO Architecture



Backbone

Feature extraction from input images



Neck

Multi-scale feature aggregation



Head

Bounding box & class prediction

→ Input → Backbone → Neck → Head → Detection

Dataset Overview



Roboflow Fruits Detection v1

Roboflow Fruits Detection Dataset



Total Images

2,974



Fruit Classes

9



Format

YOLO



Data Distribution

Training Set

2,697 (91%)

Validation Set

187 (6%)

Test Set

90 (3%)



Pre-labeled dataset for fast model fine-tuning

Annotation Format

YOLO Format

Label File Format

Each line represents one object:

```
class_id x_center y_center  
width height
```

 All values normalized to [0, 1]

Example Annotation

```
0 0.71 0.84 0.15 0.23
```



Coordinates are relative to image dimensions

Team & Roles



Beyza GÜLER

YOLOv11 Specialist & Reporting

- ✓ Trained YOLOv11 model
- ✓ Analyzed YOLOv11 metric results
- ✓ Technical documentation
- ✓ Reporting and analysis

Focus: Model evaluation & metrics analysis



Ramazan YILDIZ

AI Project Planning & Research

- ✓ Project initialization & preparation
- ✓ Research & dataset preparation
- ✓ Trained YOLOv8m & YOLOv8n models
- ✓ Model comparison analysis

Focus: Project coordination & model benchmarking



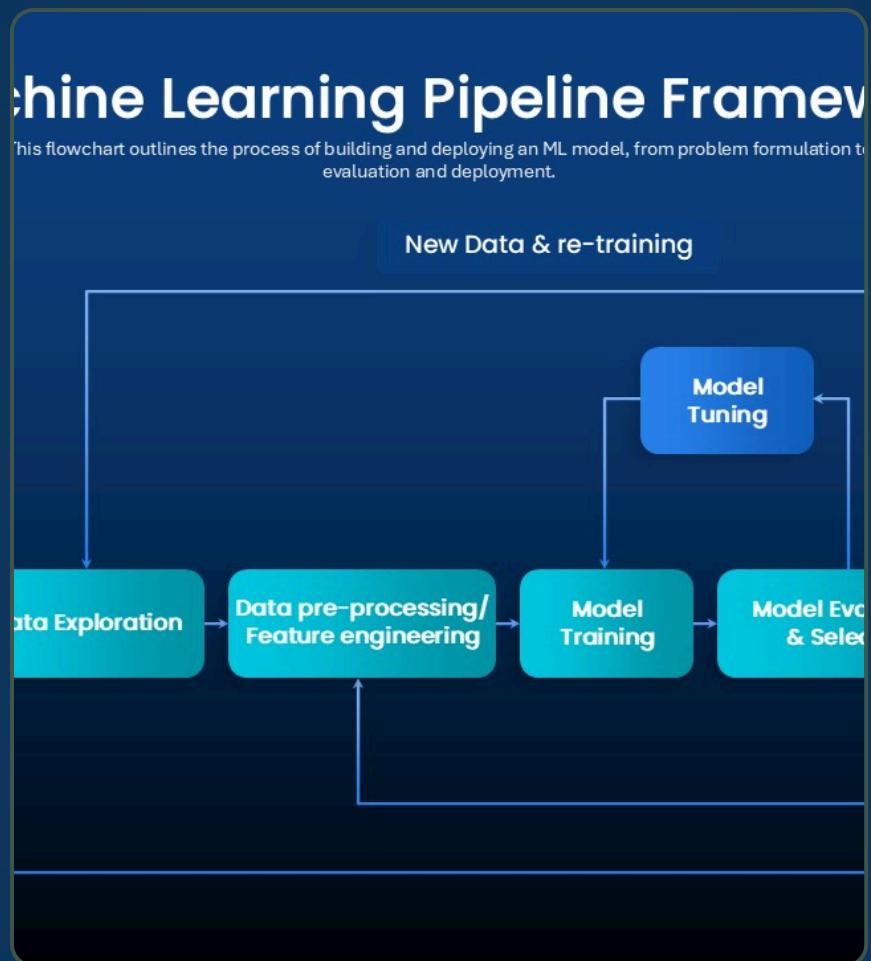
Abdelrahman MOHAMED

YOLOv5 Specialist & Web Developer

- ✓ Trained YOLOv5 model
- ✓ Analyzed YOLOv5 metrics
- ✓ Developed web interface
- ✓ OpenCV preprocessing & visualization

Focus: Web deployment & system integration

Methodology Pipeline



End-to-End Workflow



Training Setup

Environment & Configuration

Platform

- Google Colab
- GPU Acceleration

Hyperparameters

Epochs: **50**

Img Size: **640**

Batch: **16**

Device: **0 (GPU)**

💡 Consistent across all models for fair comparison

↔ Training Code Example

```
from ultralytics import YOLO  
  
model = YOLO('yolov8m.pt')  
  
results = model.train(  
    data='data.yaml',  
    epochs=50,  
    imgsz=640,  
    batch=16,  
    device=0  
)
```

💾 Best model saved as **best.pt** for deployment

YOLOv5 Results

YOLOv5 Training Results

Training Performance

Box Loss

Decreasing

Cls Loss

Converged

DFL Loss

Stable

Confusion Matrix

- ✓ Strong diagonal concentration
- ✓ High class separation accuracy
- ✓ Minimal confusion between classes

Key Metrics

mAP@0.5: ([See plot](#))

Recall: ([See plot](#))

Precision: ([See plot](#))

mAP@0.5:0.95: ([See plot](#))

YOLOv8 Results

YOLOv8 Training Results

↗ Training Progress

Train Loss

Decreasing

Val Loss

Decreasing

mAP

Increasing

.Matrix Analysis

- ✓ Strong diagonal concentration confirmed
- ✓ High accuracy across all fruit classes
- ✓ Minimal confusion between similar fruits

Performance Summary

Precision: [\(See plot\)](#)

Recall: [\(See plot\)](#)

mAP@0.5: [\(See plot\)](#)

mAP@0.5:0.95: [\(See plot\)](#)

YOLOv11 Results

YOLOv11 Training Results

↗ Training Progress

Train Loss

Rapid Drop

Val Loss

Stable

mAP

Highest

.Matrix Analysis

- ✓ Excellent class discrimination
- ✓ Strongest diagonal concentration
- ✓ Minimal false positives

★ Performance Highlights

Precision: ([See plot](#))

Recall: ([See plot](#))

mAP@0.5: ([See plot](#))

mAP@0.5:0.95: ([See plot](#))

Model Comparison

| Model | Accuracy | Speed | Strength | Weakness |
|------------|----------------|-----------------|-------------------|-------------------|
| YOLOv5m ⭐ | (From results) | Fastest | Fastest inference | Moderate accuracy |
| YOLOv8m ⭐ | (From results) | Balanced | Good performance | Moderate speed |
| YOLOv11m ⭐ | Highest | Moderate | Best accuracy | Slower than v5 |

-Key Findings

- ✓ YOLOv11m achieved **highest detection accuracy**
- ✓ YOLOv5m provided **fastest inference time**

Real-Time Performance

Achieved FPS:

~30

Measured on development machine

Models tested under
identical conditions

Deployment Architecture

System Pipeline

1 Input Sources
Webcam, Video, Images

2 Preprocessing
OpenCV capture & resize

3 Inference
YOLO model detection

4 Postprocessing
Draw boxes & labels

5 UI Visualization
Streamlit web app

6 Analytics Logging
Detection summary

⌚ End-to-end latency: ~32ms/frame

Technology Stack



Python



YOLO



PyTorch



OpenCV



Streamlit



WebRTC

User Interface

⬆ Overview Page

📊 Analytics Dashboard

🎥 Live Detection

🌐 Model Selection

✅ Real-time performance: ~30 FPS

☁️ Browser-based deployment

Web Application - Overview & Live Detection



Streamlit Web Interface

Overview & Live Detection Pages

Overview Page Features

Project Introduction

YOLOv8 architecture overview

Technology Stack

Python, YOLO, OpenCV, Streamlit

Model Details

YOLOv8n, 80+ classes, ~30 FPS

System Status

Active & operational

Live Detection Capabilities



Bounding Boxes



Class Labels



Confidence

Real-Time Performance

~30 FPS

Browser-Based

Easy access

Analytics Dashboard



Real-Time Analytics

Detection statistics & visualization

II. Detection Statistics

Total Detections

1,250

Avg Confidence

84.5%

Unique Classes

18

Processing Time

32ms

✓ Detection Frequency

- Real-time object distribution
- Bar chart visualization
- Most frequent objects

📊 Top Detections

| | |
|--------|-----------------|
| Person | High |
| Car | Medium |
| Fruits | Variable |

⚡ Detection Details

Per-class statistics & logs

Challenges & Limitations

⚠ Key Challenges



Accuracy vs Speed Tradeoff

Larger models more accurate but slower



Lighting & Occlusion

Detection quality varies with conditions



Small Objects & Crowds

Harder to detect in complex scenes

ℹ Current Limitations



Limited Dataset

Only 9 fruit classes, need diversity



Hardware Constraints

GPU availability & compute resources



Generalization Issues

Performance drops in unseen environments



Lab Constraints

Webcam permissions & driver compatibility

Backup Plan & Risk Mitigation

Model Issues

- Use stable YOLOv8 or switch to YOLOv5
- Alternative frameworks available

Training Instability

- Reduce classes or data subset
- Adjust hyperparameters (LR, batch, epochs)

Data Problems

- Use additional public datasets
- Collect new training samples

Compute Load

- Migrate to cloud GPU platform
- Use lighter nano model variant

Web App Issues

- Alternative UI frameworks (Gradio, Dash)
- Core detection remains functional

Timeline Delays

- Overlap less critical project steps
- Extend schedule maintaining objectives

Mitigation Strategy

Proactive backup measures ensure project continuity

Project Roadmap



Development Timeline

Current & Next Phases



Phase 1: Completed

Current Status ✓

- ✓ Dataset sampling completed
- ✓ Real-time detection working
- ✓ Prototype system deployed
- ✓ Web interface established



Phase 2: Next Steps

Planned Development

- ➡ Image annotation & labeling
- ➡ Quantitative evaluation

➡ Model fine-tuning (bigger dataset)

➡ Performance optimization



Ready for Phase 2 Development

Continuous Improvement

Conclusion



Working Prototype

Real-time object detection successfully implemented and deployed



System Integration

YOLO + OpenCV + Streamlit web application fully functional



Performance Achievement

Stable real-time performance at **~30 FPS** achieved



Comprehensive Evaluation

- ✓ Benchmarked 3 YOLO architectures
- ✓ Produced training artifacts
- ✓ Generated evaluation metrics



Browser Deployment

- 💻 Accessible via web browser
- 📷 Real-time webcam detection
- 📊 Analytics dashboard included



Strong foundation for future enhancements

Future Work & Thank You

❖ Enhanced Evaluation

- ✓ Per-class AP analysis
- ✓ Model quantization
- ✓ Latency profiling
- ✓ Rigorous benchmarking

□ Deployment Expansion

- ✓ Mobile app deployment
- ✓ On-device inference
- ✓ Edge computing optimization

↑ UI Improvements

- ✓ Model switcher
- ✓ Save results feature
- ✓ Threshold controls
- ✓ Session logs

▪ Dataset Expansion

- ✓ Larger custom dataset training
- ✓ More object classes
- ✓ Varied environmental conditions

▣ Advanced Features

- ✓ 3D localization & depth estimation
- ✓ Multi-object tracking
- ✓ Instance segmentation



Thank You!

We appreciate your attention and feedback

❑ Questions & Answers

